

Advanced Modeling

Cam Smithers

2026-07-05

Necessary Packages

```
library(tidyverse)
library(rsample)
library(caret)
library(vip)
library(ROCR)
library(nnet)
```

Necessary Data Sets

```
setwd('/Users/camsmithers/Desktop/Camalytics/NBA')
teamvisdata <- readRDS('Data-NBA/Current/teamvisdata.rds')
teamvisdata_fixed <- na.omit(teamvisdata) %>%
  mutate(
    team = as.factor(team),
    opponent = as.factor(opponent),
    season = as.factor(season))
```

Logistic Regression Model Testing: Binary Game Outcome

```
winbinary_split <- initial_split(teamvisdata_fixed, prop=.75, strata='twoleveloc')
winbinary_train <- training(winbinary_split)
winbinary_test <- testing(winbinary_split)
```

For the models I started by adding all the variables that I wanted and deciding which ones I wanted to remove to see how the accuracy changes.

```
winbinary_1 <- train(
  twoleveloc ~ roll_offrating + roll_defrating + team + homeaway +
    vorp_bpa + roll_three_par + roll_ftar + roll_ast_pct,
  data=teamvisdata_fixed,
  method='glm',
  family='binomial',
  trControl=trainControl(method='cv', number=10))
winbinary_2 <- train(
  twoleveloc ~ roll_offrating + roll_defrating,
  data=teamvisdata_fixed,
  method='glm',
  family='binomial',
  trControl=trainControl(method='cv', number=10))
```

```
winbinary_3 <- train(
  twoleveloc ~ roll_offrating + opponent,
  data=teamvisdata_fixed,
  method='glm',
  family='binomial',
  trControl=trainControl(method='cv', number=10))
winbinary_4 <- train(
  twoleveloc ~ homeaway,
  data=teamvisdata_fixed,
  method='glm',
  family='binomial',
  trControl=trainControl(method='cv', number=10))
```

This will give the me summary statistics of the models' accuracy.

```
summary(
  resamples(
    list(
      model1=winbinary_1,
      model2=winbinary_2,
      model3=winbinary_3,
      model4=winbinary_4
    )
  )
)$statistics$Accuracy
```

##		Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	model1	0.6478029	0.6555223	0.6618964	0.6614976	0.6678310	0.6719894	0
##	model2	0.6420492	0.6505242	0.6533599	0.6583047	0.6689953	0.6779774	0
##	model3	0.6200931	0.6210912	0.6269547	0.6270290	0.6305722	0.6420492	0
##	model4	0.5199734	0.5430183	0.5582169	0.5545617	0.5680306	0.5741850	0

ROC Curve for Logistic Regression Models

Creating a plot to show the False Positive vs. True Positive Accuracy.

```
model1_prob <- predict(winbinary_1, winbinary_train, type='prob')$Win
perf1 <- prediction(model1_prob, winbinary_train$twoleveloc) %>%
  performance(measure='tpr', x.measure='fpr')

model2_prob <- predict(winbinary_2, winbinary_train, type='prob')$Win
perf2 <- prediction(model2_prob, winbinary_train$twoleveloc) %>%
  performance(measure='tpr', x.measure='fpr')

model3_prob <- predict(winbinary_3, winbinary_train, type='prob')$Win
perf3 <- prediction(model3_prob, winbinary_train$twoleveloc) %>%
  performance(measure='tpr', x.measure='fpr')

model4_prob <- predict(winbinary_4, winbinary_train, type='prob')$Win
perf4 <- prediction(model4_prob, winbinary_train$twoleveloc) %>%
  performance(measure='tpr', x.measure='fpr')

plot(perf1, col='cornflowerblue', lty=1)
plot(perf2, add=TRUE, col='magenta2', lty=2)
plot(perf3, add=TRUE, col='gold2', lty=2)
```

```

plot(perf4, add=TRUE, col='forestgreen', lty=2)
legend(0.6, 0.3, legend=c('Full Model',
                          'Net Rating',
                          'Offensive Rating w/ Opponent',
                          'Defensive Rating w/ Team'),
      col=c('cornflowerblue', 'magenta2', 'gold2', 'forestgreen'),
      lty=2:1, cex=0.6)

```

